Kubernetes Empowerer via API (KEA)

Task 3

Project statement

The product is a platform for deploying, managing, and scaling machine learning models in a production environment. It's primary purpose is to provide a flexible and secure environment for automating ML processes, including model versioning, request routing, and monitoring. The system integrates with Kubernetes, supports model containerization. The product is designed for developers, ML engineers, DevOps teams, and enterprises that require a stable, scalable, and resilient infrastructure for their ML projects.

Team K8C: Daniel Tsurkan; Gadji Dandamaev; Tsaturyan Konstantin; Smolkin Mikhail

Project repo: https://github.com/fanglores/Advanced-Software-Design

This report: here

Roles

ML Engineer

Description: This role joins professionals involved in the development, deployment, and monitoring of ML models. They want to simplify the deployment process, automate API documentation, and ensure efficient request validation and caching, ultimately enhancing their workflow and model performance.

API Consumer

Description: This role includes all users interacting with APIs to integrate ML models into their applications. They want to access reliable and well-documented APIs, enabling seamless integration of ML models into their business applications and ensuring optimal performance and usability.

Personas

ML Engineer (Maria, 32 years old)

Goals:

- Deploy and version ML models in Kubernetes.
- Automatic API documentation and request validation.
- Flexibility for different ML frameworks.

Pain points:

- Manual API documentation.
- Difficulties in monitoring model performance.

Backend Developer (Alexander, 28 years old)

Goals:

- Use automatic OpenAPI schema generation.
- Easily add API endpoints with request validation and security.

Pain points:

- Manual API documentation.
- Challenges with integrating authorization and managing access control.

Personas

API Consumer (Sergey, 30 years old)

Goals:
- Get documentation for quick access to ML models.
- Work with reliable and validated APIs.

Pain points:
- Incomplete or outdated documentation.
- API instability and delays.

Corporate Client (Yandex, Sber)

Goals:
- Scalable and secure deployment of ML models.
- Integration of the API gateway into existing infrastructure.

Pain points:
- Challenges with integration and corporate standards.

DFD (Level 0)

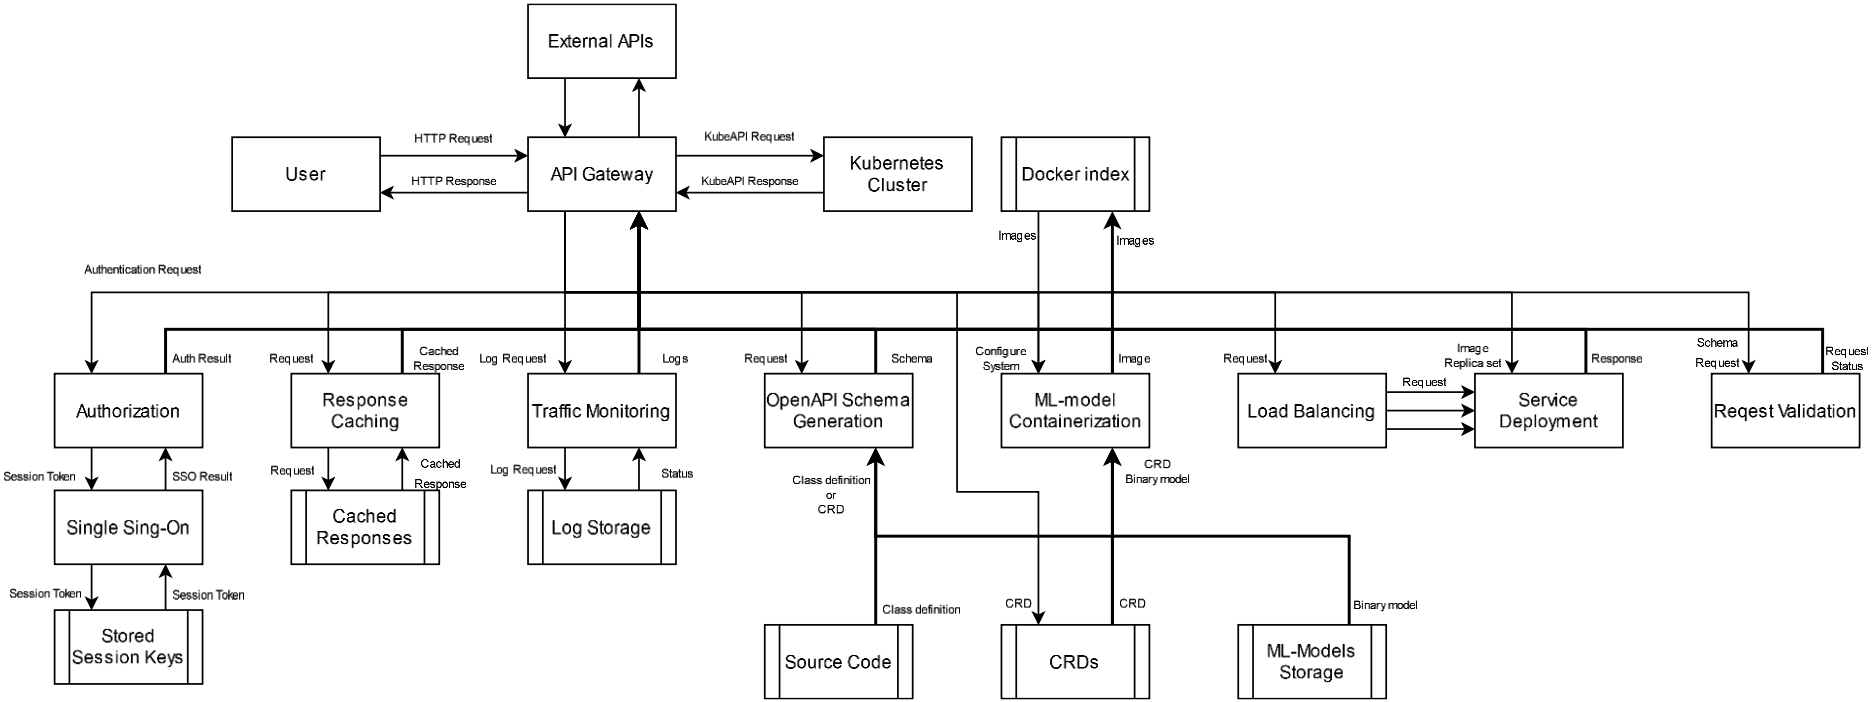
External APIs
User
HTTP Request
HTTP Response
API Gateway
KubeAPI Request
KubeAPI Response
Kubernetes Cluster
Docker index
Images
Images
Authentication Request
Auth Result
Request
Cached Response
Log Request
Logs
Request
Schema
Configure System
Image
Request
Image Replica set
Response
Schema Request
Request Status
Authorization
Response Caching
Traffic Monitoring
OpenAPI Schema Generation
ML-model Containerization
Load Balancing
Request
Service Deployment
Reqest Validation
Session Token
SSO Result
Request
Cached Response
Log Request
Status
Class definition or CRD
CRD Binary model
Single Sing-On
Cached Responses
Log Storage
Session Token
Session Token
Class definition
CRD
CRD
Binary model
Stored Session Keys
Source Code
CRDs
ML-Models Storage

Story map

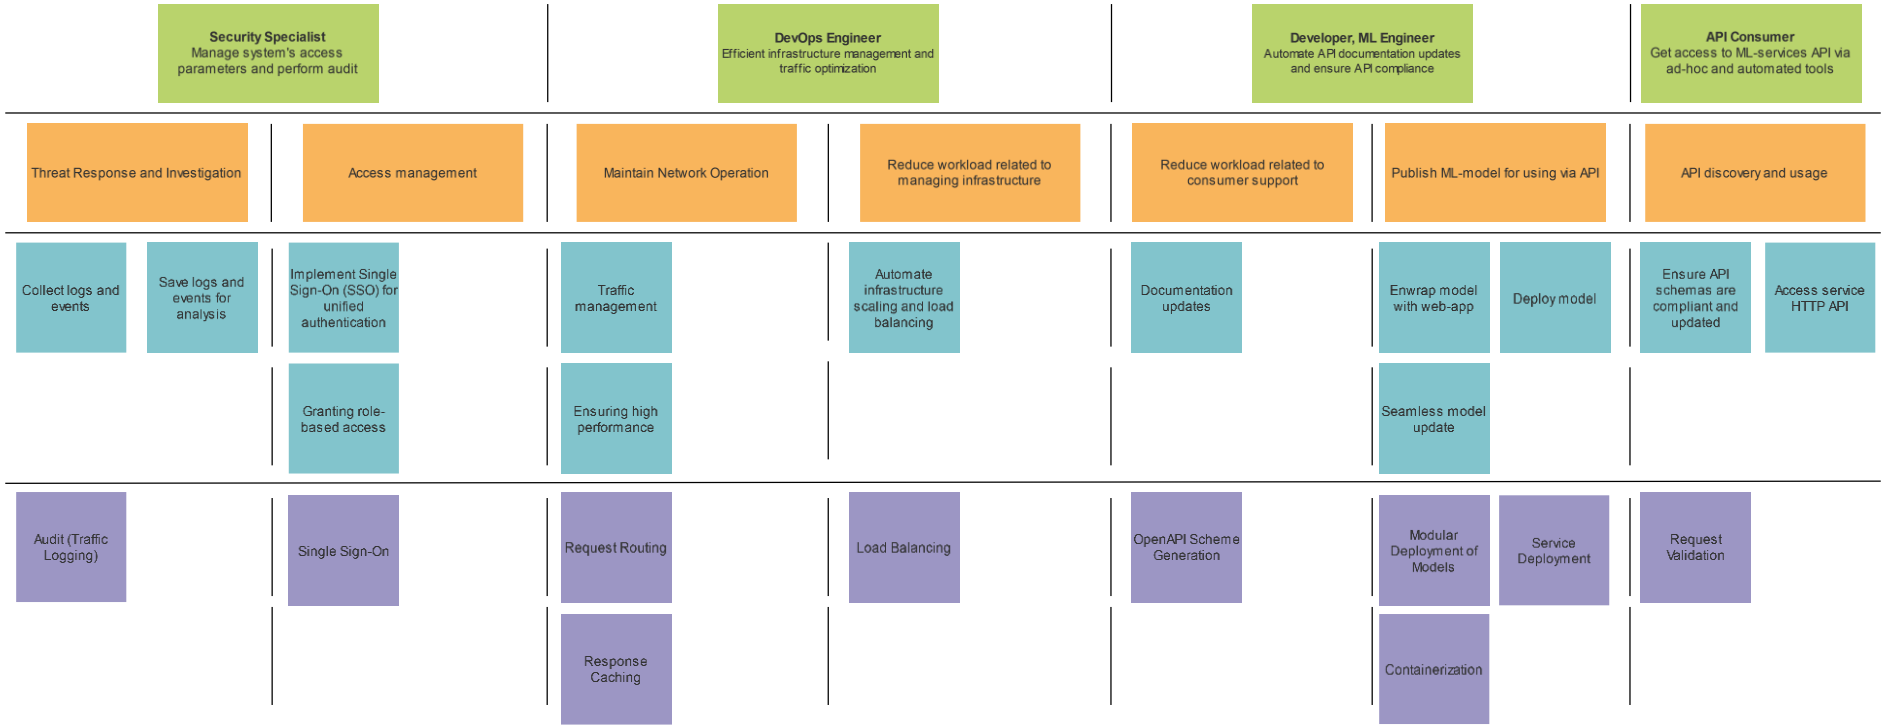
Security Specialist
Manage system's access parameters and perform audit
DevOps Engineer
Efficient infrastructure management and traffic optimization
Developer, ML Engineer
Automate API documentation updates and ensure API compliance
API Consumer
Get access to ML-services API via ad-hoc and automated tools
Threat Response and Investigation
Access management
Maintain Network Operation
Reduce workload related to managing infrastructure
Reduce workload related to consumer support
Publish ML-model for using via API
API discovery and usage
Collect logs and events
Save logs and events for analysis
Implement Single Sign-On (SSO) for unified authentication
Granting role-based access
Traffic management
Ensuring high performance
Automate infrastructure scaling and load balancing
Documentation updates
Enwrap model with web-app
Seamless model update
Deploy model
Ensure API schemas are compliant and updated
Access service HTTP API
Audit (Traffic Logging)
Single Sign-On
Request Routing
Response Caching
Load Balancing
OpenAPI Scheme Generation
Modular Deployment of Models
Containerization
Service Deployment
Request Validation