

KEA

(Kubernetes Empowerer to API)

Requirements and analysis model

Product description

The product is a platform for deploying, managing, and scaling machine learning models in production. It offers a secure, flexible environment for automating ML tasks like model versioning, routing, and monitoring. With Kubernetes integration and containerization support, it's designed for developers, ML engineers, and enterprises needing scalable, resilient ML infrastructure.

Team K8C: Tsurkan Daniel; Dandamaev Gadji; Tsaturyan Konstantin; Smolkin Mikhail

Project repo: <https://github.com/fanglores/Advanced-Software-Design>

This report: [https://github.com/fanglores/Advanced-Software-Design/blob/master/Practice%20Tasks/Final_Task/K8C_FinalTask1_\(Task7\).pdf](https://github.com/fanglores/Advanced-Software-Design/blob/master/Practice%20Tasks/Final_Task/K8C_FinalTask1_(Task7).pdf)

Roles

ML Engineer

Description: This role joins professionals involved in the development, deployment, and monitoring of ML models. They want to simplify the deployment process, automate API documentation, and ensure efficient request validation and caching, ultimately enhancing their workflow and model performance.

API Consumer

Description: This role includes all users interacting with APIs to integrate ML models into their applications. They want to access reliable and well-documented APIs, enabling seamless integration of ML models into their business applications and ensuring optimal performance and usability.

Personas

ML Engineer (Maria, 32 years old)

Goals:

- Deploy and version ML models in Kubernetes.
- Automatic API documentation and request validation.
- Flexibility for different ML frameworks.

Pain points:

- Manual API documentation.
- Difficulties in monitoring model performance.

Backend Developer (Alexander, 28 years old)

Goals:

- Use automatic OpenAPI schema generation.
- Easily add API endpoints with request validation and security.

Pain points:

- Manual API documentation.
- Challenges with integrating authorization and managing access control.

Personas

API Consumer (Sergey, 30 years old)

Goals:

- Get documentation for quick access to ML models.
- Work with reliable and validated APIs.

Pain points:

- Incomplete or outdated documentation.
- API instability and delays.

Corporate Client (Yandex, Sber)

Goals:

- Scalable and secure deployment of ML models.
- Integration of the API gateway into existing infrastructure.

Pain points:

- Challenges with integration and corporate standards.

Story map

Security Specialist
Manage system's access parameters and perform audit

DevOps Engineer
Efficient infrastructure management and traffic optimization

Developer, ML Engineer
Automate API documentation updates and ensure API compliance

API Consumer
Get access to ML-services API via ad-hoc and automated tools

Threat Response and Investigation

Access management

Maintain Network Operation

Reduce workload related to managing infrastructure

Reduce workload related to consumer support

Publish ML-model for using via API

API discovery and usage

Collect logs and events

Save logs and events for analysis

Implement Single Sign-On (SSO) for unified authentication

Traffic management

Automate infrastructure scaling and load balancing

Documentation updates

Enwrap model with web-app

Deploy model

Ensure API schemas are compliant and updated

Access service HTTP API

Granting role-based access

Ensuring high performance

Seamless model update

Audit (Traffic Logging)

Single Sign-On

Request Routing

Load Balancing

OpenAPI Scheme Generation

Modular Deployment of Models

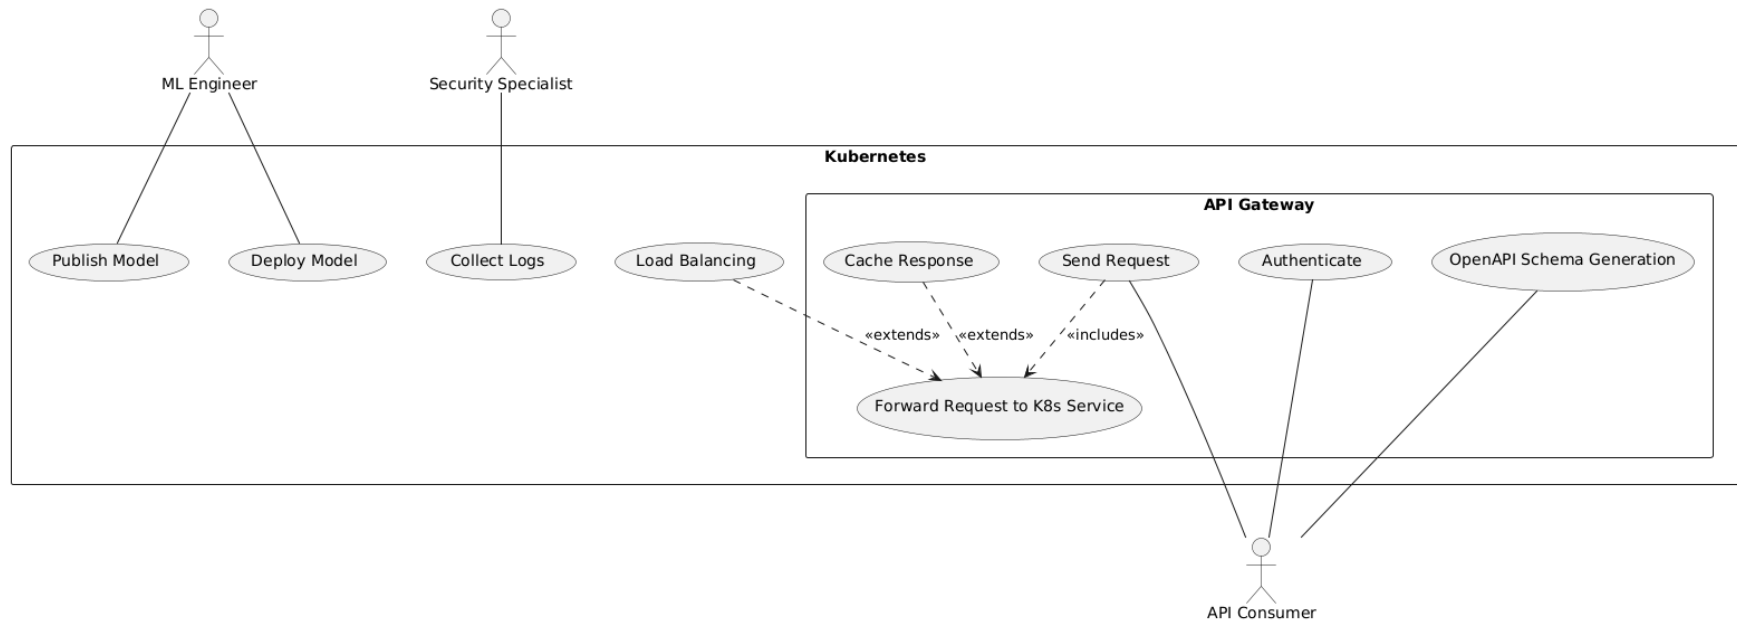
Service Deployment

Request Validation

Response Caching

Containerization

Use case diagram



Interaction analysis

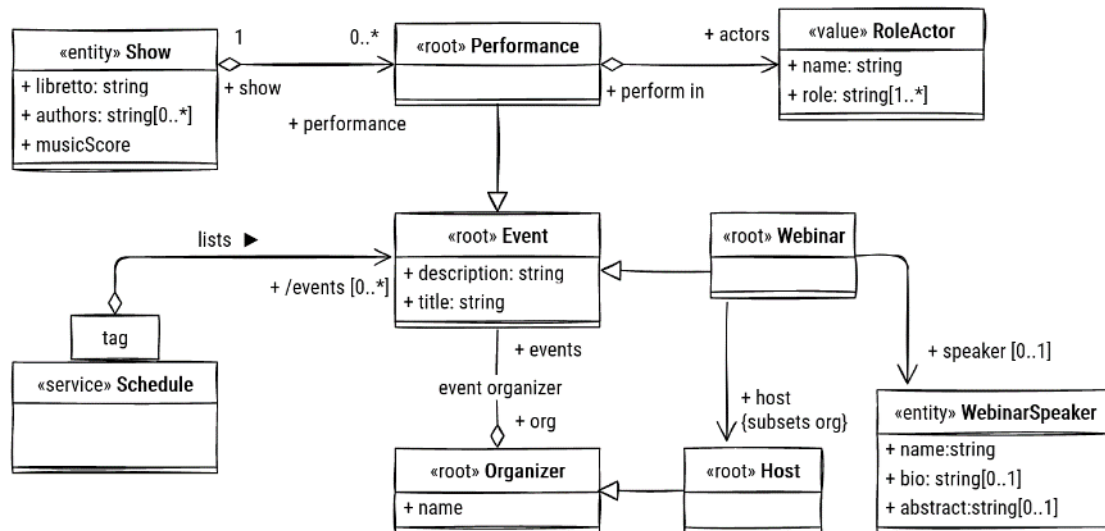
Show a cooperation on a diagram along with the use case

OR. Show as a table, columns: use case - cooperation name - used roles - candidate classes that play them

Final class diagram

The final class model should be consistent (well-formed) with interactions and use cases / stories.

Check that DDD stereotypes are set



Detailed behaviour

Use an activity or state diagram to describe behavior of **one** of the dynamic classifiers in your model

Add the diagram here and explain how it works precisely in detail

Check that the overall model remains well-formed

Note: too small models may not be enough to demonstrate your qualification

Repository structure


Repository structure

```
.
├── product_img.jpg
├── README.md
├── General/
│   ├── Domain_description_en.md
│   ├── Task_description_en.md
│   ├── DFD0/
│   ├── StoryMap/
│   └── UseCases/
├── Practice Tasks/
│   ├── Final_Task/
│   ├── Task_1/
│   ├── Task_2/
│   ├── Task_3/
│   ├── Task_4/
│   ├── Task_5/
│   ├── Task_6/
│   └── Task_7/
```

Tools Used:

- Github
- Drawio
- Planttext

Team and roles




Roles, Personas, Story Map	Use cases, DDD	UML, DFD, Repository Management	Domain analysis, CRC Cards, UML
Tsurkan Daniel Tg: @crazy_deyzi	Dandamaev Gadji Tg: @dandamaev	Tsaturyan Konstantin Tg: @fanglores	Smolkin Mikhail Tg: @m0hnatik

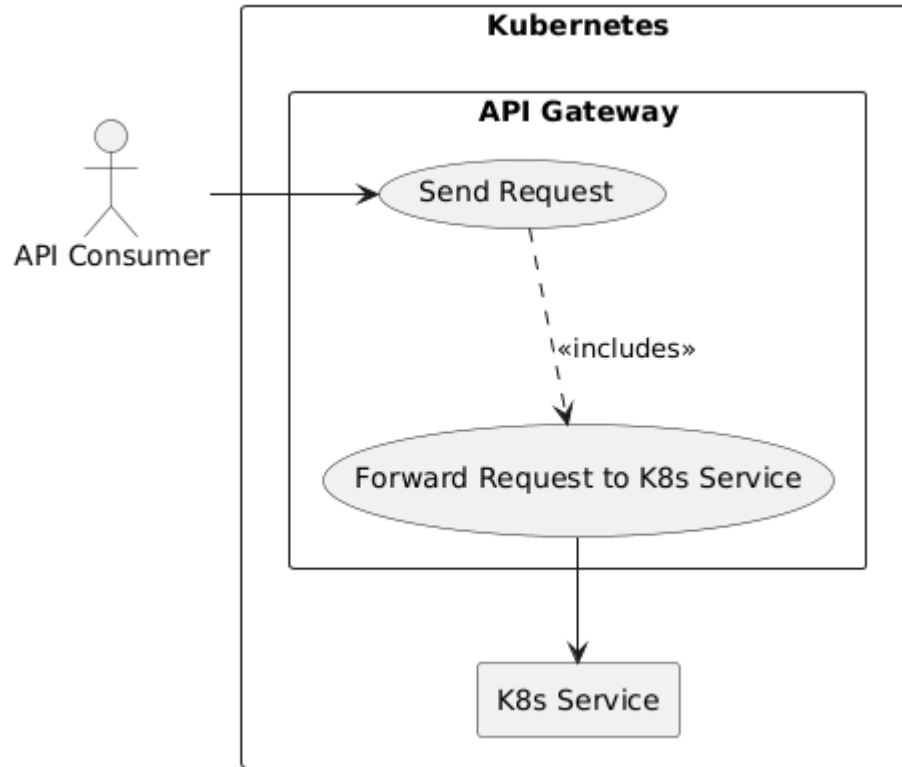
Thanks for attention!

Now we are ready to answer your questions!

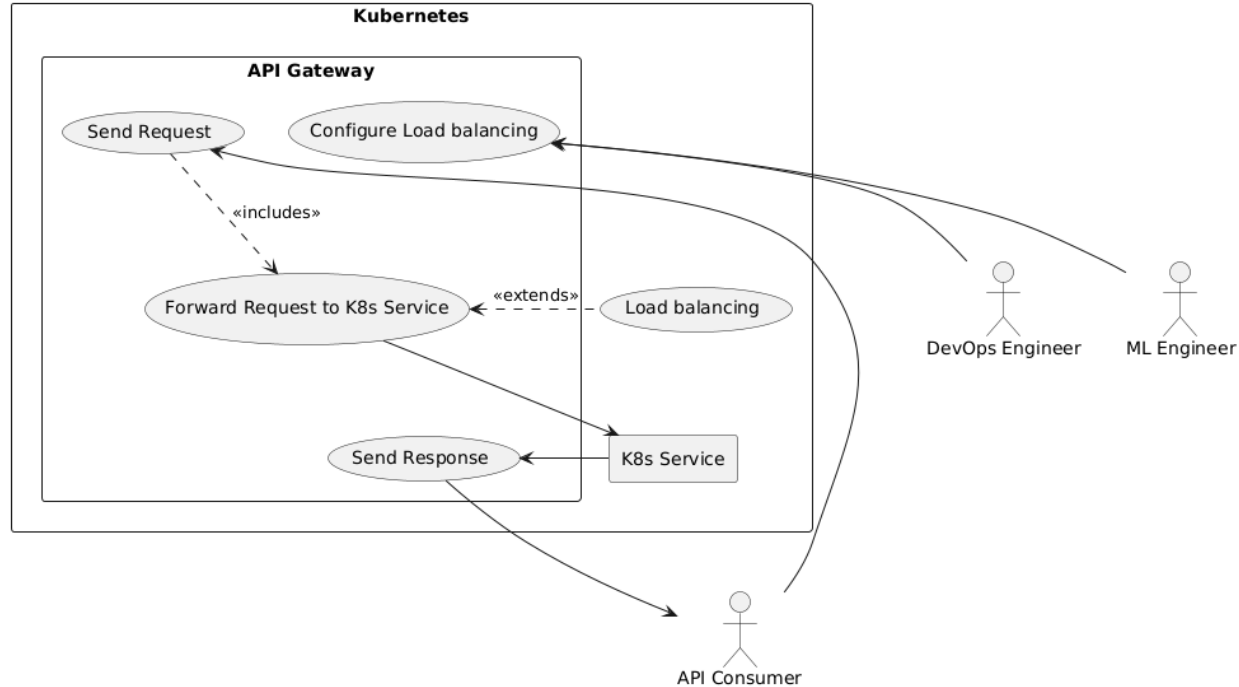
DFD (Level 0)




Use Cases: Requests routing



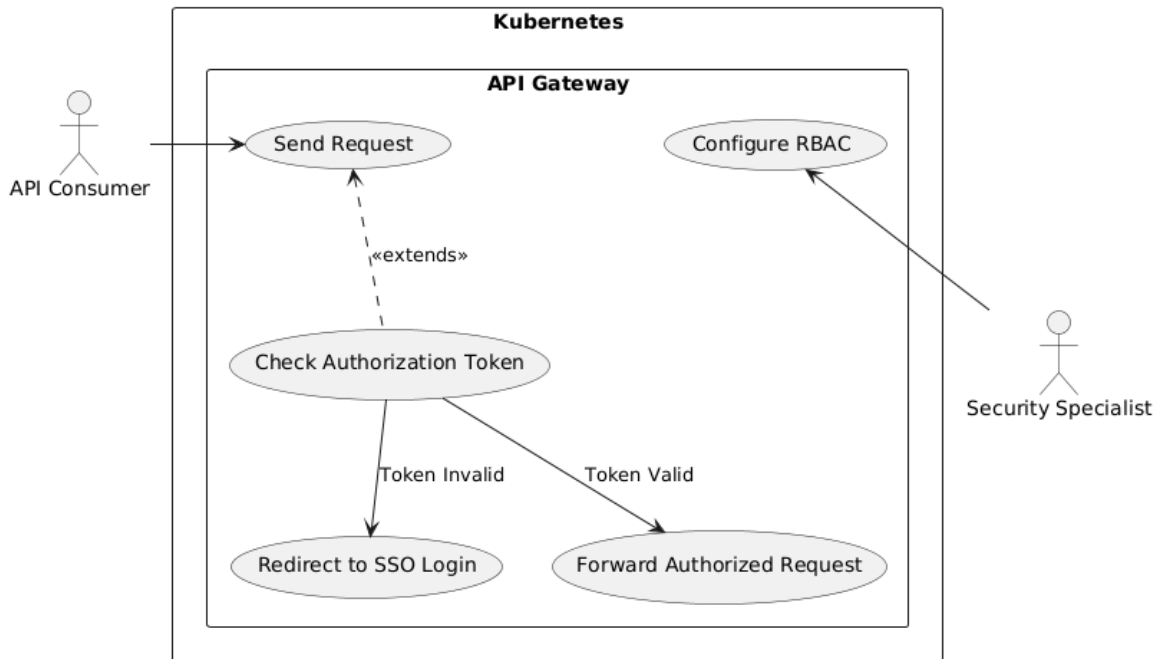
Use Cases: Load balancing



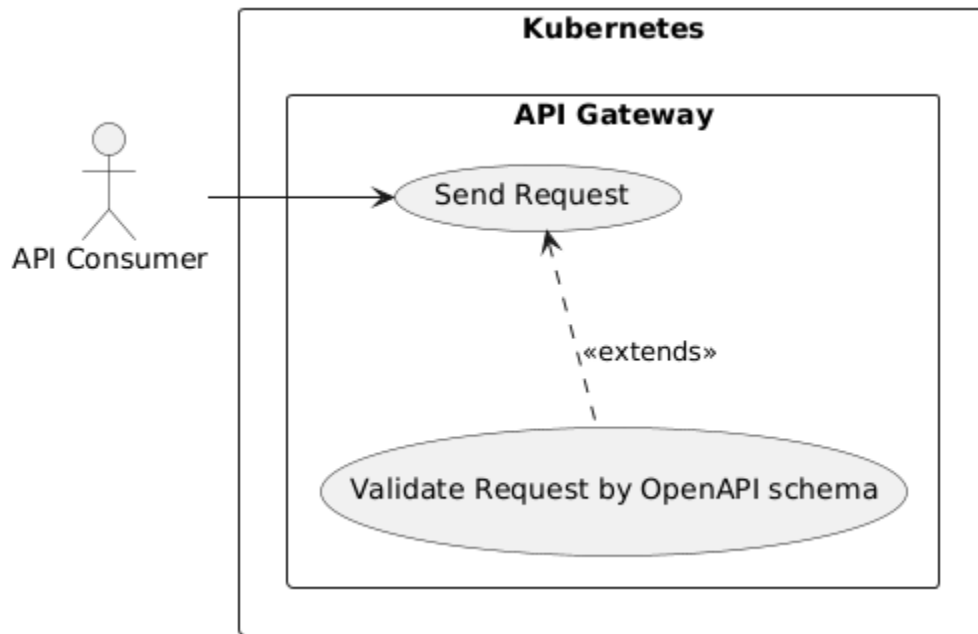
Use Cases: Audit and Logging




Use Cases: Authorization (SSO)



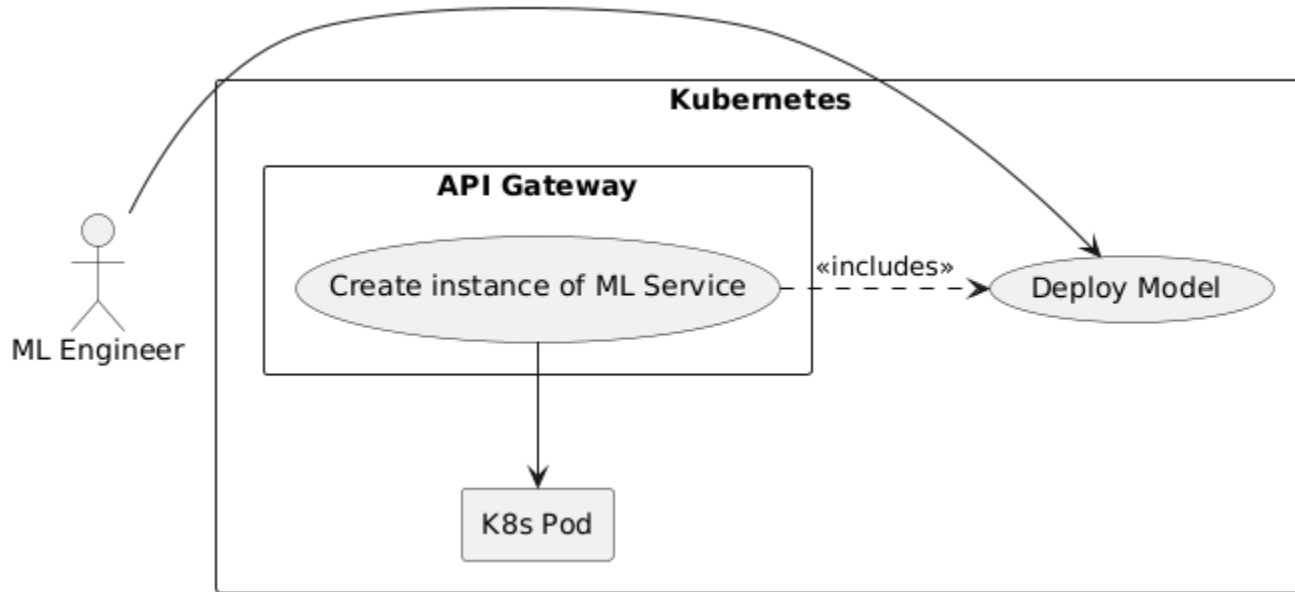
Use Cases: Request Validation




Use Cases: Responses Caching



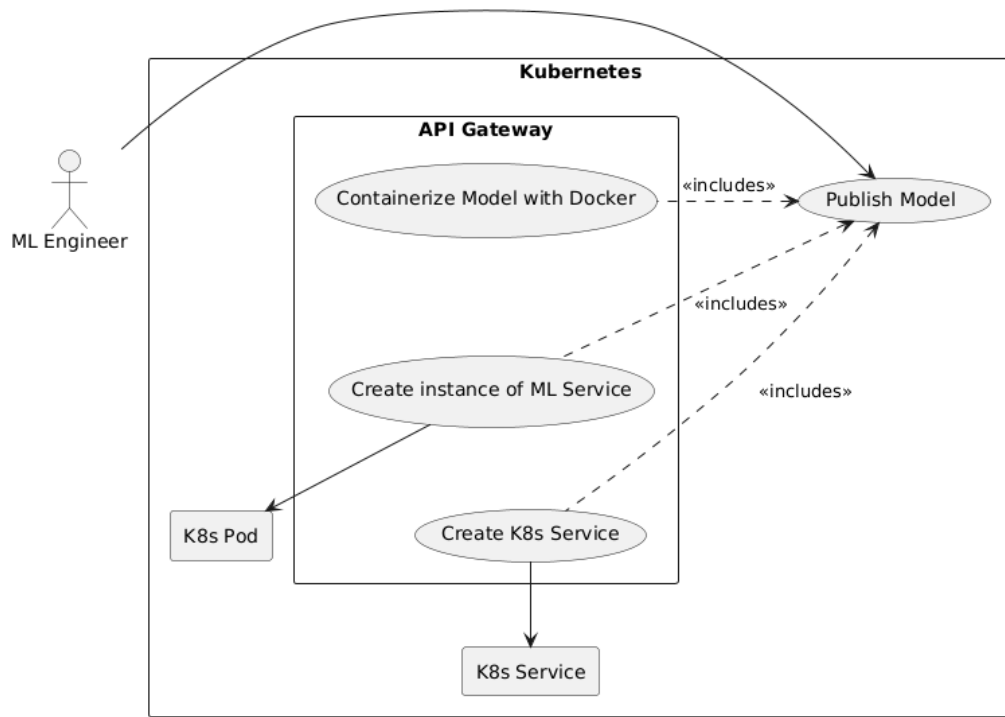
Use Cases: Modular Deployment of Models



Use Cases: Containerization



Use Cases: Service Deployment



Use Cases: Auto-Documentation

